

# **Analysing MQM error profiles and metric correlations in MT systems**

Multilingual NLP – Lab 3

Philippos Triantafyllou

January 14, 2026

## **Table of contents**

<b>1 Data loading</b>	<b>2</b>
<b>2 Building and visualizing error profiles</b>	<b>3</b>
<b>3 Clustering systems by error profiles</b>	<b>5</b>
<b>4 MQM and Automatic Metrics</b>	<b>6</b>
<b>5 Linguistic factors and correlation</b>	<b>15</b>

# 1 Data loading

```

import pandas as pd

pd.set_option("display.max_columns", 8)
pd.set_option("display.max_colwidth", 20)
pd.set_option("display.width", 80)

# Load the TSV file
df = pd.read_csv("data/mqm_2022_ende.tsv", sep="\t")

# # Select required columns
cols = ["system_id", "doc_segment_id", "category", "severity", "source", "candidate", "reference"]
df = df[cols]
df = df.rename(columns={'system_id': 'system', 'candidate': 'target', 'doc_segment_id': 'segment'})

display(df)
print(df.shape)

```

	system	segment	category	severity	source	target
0	JDExploreAcademy	1	style/awkward	minor	Avalanche at Was...	Lawine in Skig...
1	JDExploreAcademy	1	style/awkward	minor	Avalanche at Was...	Lawine in Skig...
2	JDExploreAcademy	1	no-error	no-error	Avalanche at Was...	Lawine in Skig...
3	JDExploreAcademy	1	accuracy/mistrans...	minor	Avalanche at Was...	Lawine in Skig...
4	Lan-Bridge	1	style/awkward	minor	Avalanche at Was...	Lawine im Skig...
...	...	...	...	...	...	...
78554	comet_bestmbr	10	no-error	no-error	I shall never fo...	Ich werde dies...
78555	refB	10	fluency/punctuation	minor	I shall never fo...	<v>Ich</v>
78556	refB	10	fluency/punctuation	minor	I shall never fo...	Ich werde dies...
78557	refB	10	no-error	no-error	I shall never fo...	Ich werde dies...
78558	refB	10	no-error	no-error	I shall never fo...	Ich werde dies...

(78559, 7)

Doing some cleaning.

```

# Drop rows with missing annotations in key fields
df = df.dropna(subset=["system", "segment", "category", "severity", "source", "target",
                     "reference"]).reset_index(drop=True)
display(df.shape)

```

(78559, 7)

```

# View the unique category and severity labels
print("Unique Categories:", df["category"].unique())
print("Unique Severities:", df["severity"].unique())

```

```

Unique Categories: ['style/awkward' 'no-error' 'accuracy/mistranslation'
 'accuracy/untranslated' 'fluency/grammar' 'accuracy/addition'
 'accuracy/omission' 'terminology/inappropriate' 'locale_convention/time'
 'fluency/register' 'terminology/inconsistent' 'fluency/spelling'
 'fluency/punctuation' 'fluency/display' 'fluency/inconsistency'
 'locale_convention/currency' 'source_error' 'other' 'non_translation']

```

```
'locale_convention/date' 'locale_convention/name']
Unique Severities: ['minor' 'no-error' 'major']

# Build a major category variable
df["major_category"] = df["category"].apply(lambda x: x.split("/")[0].split(".")[0].strip() if pd.notnull(x)
                                         else x)
df["major_category"].unique()

array(['style', 'no-error', 'accuracy', 'fluency', 'terminology',
       'locale_convention', 'source_error', 'other', 'non_translation'],
      dtype=object)
```

The data is now ready to work with.

## 2 Building and visualizing error profiles

First we build for each system its error profile as a vector of error counts per major category.

```
df_err = df[df["major_category"] != "no-error"].copy()

# Count errors per system and category
error_counts = (df_err.groupby(["system", "major_category"]).size().unstack(fill_value=0))

display(error_counts)
```

major_category system	accuracy	fluency	locale_convention	non_translation	other	source_error
JDExploreAcademy	957	874	4	0	1	14
Lan-Bridge	1233	1362	5	0	4	19
M2M100_1.2B-B4	2681	2054	2	3	9	17
Online-A	1151	1468	4	1	4	27
Online-B	940	967	5	0	8	24
Online-G	1177	1516	16	1	20	18
Online-W	673	1082	9	0	2	15
Online-Y	1256	1312	7	0	14	19
OpenNMT	1657	921	4	0	5	18
PROMT	1564	1399	3	0	1	24
QUARTZ_TuneReranking	1287	1482	9	0	10	13
bleu_bestmbr	920	988	5	0	8	14
bleurt_bestmbr	976	986	4	0	4	17
comet_bestmbr	1051	1092	4	0	1	21
refB	772	614	11	0	0	18

We can see that most errors are either accuracy or fluency errors.

```
# count target-side tokens
df_err["target_tokens"] = df_err["target"].str.split().str.len()

tokens_per_system = (df_err.groupby("system")["target_tokens"].sum())

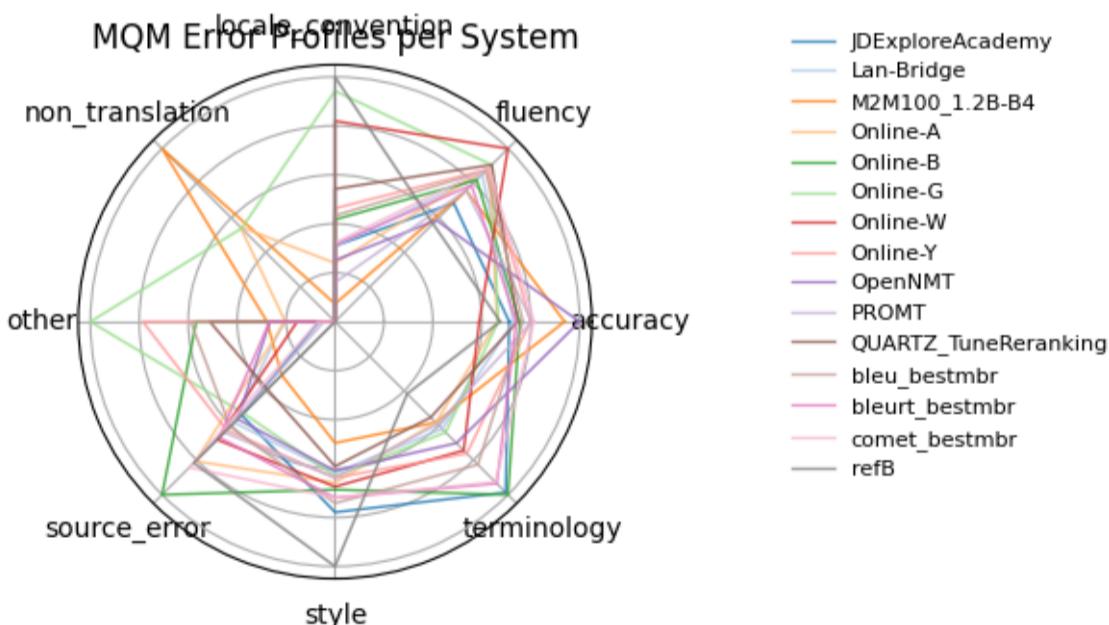
error_freq_per_token = error_counts.div(tokens_per_system, axis=0)
```

We can normalize by total number of tokens to get error rates.

```
display(error_freq_per_token)
```

major_category system	accuracy	fluency	locale_convention	non_translation	other	source_error
JDExploreAcademy	0.017037	0.015560	0.000071	0.000000	0.000018	0.000249
Lan-Bridge	0.017461	0.019287	0.000071	0.000000	0.000057	0.000269
M2M100_1.2B-B4	0.022451	0.017200	0.000017	0.000025	0.000075	0.000142
Online-A	0.015857	0.020225	0.000055	0.000014	0.000055	0.000372
Online-B	0.018068	0.018587	0.000096	0.000000	0.000154	0.000461
Online-G	0.015959	0.020555	0.000217	0.000014	0.000271	0.000244
Online-W	0.014108	0.022682	0.000189	0.000000	0.000042	0.000314
Online-Y	0.019116	0.019969	0.000107	0.000000	0.000213	0.000289
OpenNMT	0.023920	0.013295	0.000058	0.000000	0.000072	0.000260
PROMT	0.019253	0.017222	0.000037	0.000000	0.000012	0.000295
QUARTZ_TuneReranking	0.017857	0.020562	0.000125	0.000000	0.000139	0.000180
bleu_bestmbr	0.018387	0.019746	0.000100	0.000000	0.000160	0.000280
bleurt_bestmbr	0.017721	0.017902	0.000073	0.000000	0.000073	0.000309
comet_bestmbr	0.019453	0.020212	0.000074	0.000000	0.000019	0.000389
refB	0.016153	0.012847	0.000230	0.000000	0.000000	0.000377

Then we visualize using a radar plot.



Overall, systems have relatively similar profiles, especially when it comes to frequent error types such as accuracy, fluency, terminology and style. However, there is a gradual increase in these metrics (we can distinguish the radius of individual systems). Less frequent error types

such as source errors and non-translation show more distinct system profiles, and already we can notice some system particularities. For example M2M has a lot of non-translation errors.

### 3 Clustering systems by error profiles

```
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA

error_frequency_dataframe = error_freq_per_token
system_names = error_frequency_dataframe.index
feature_matrix = error_frequency_dataframe.values

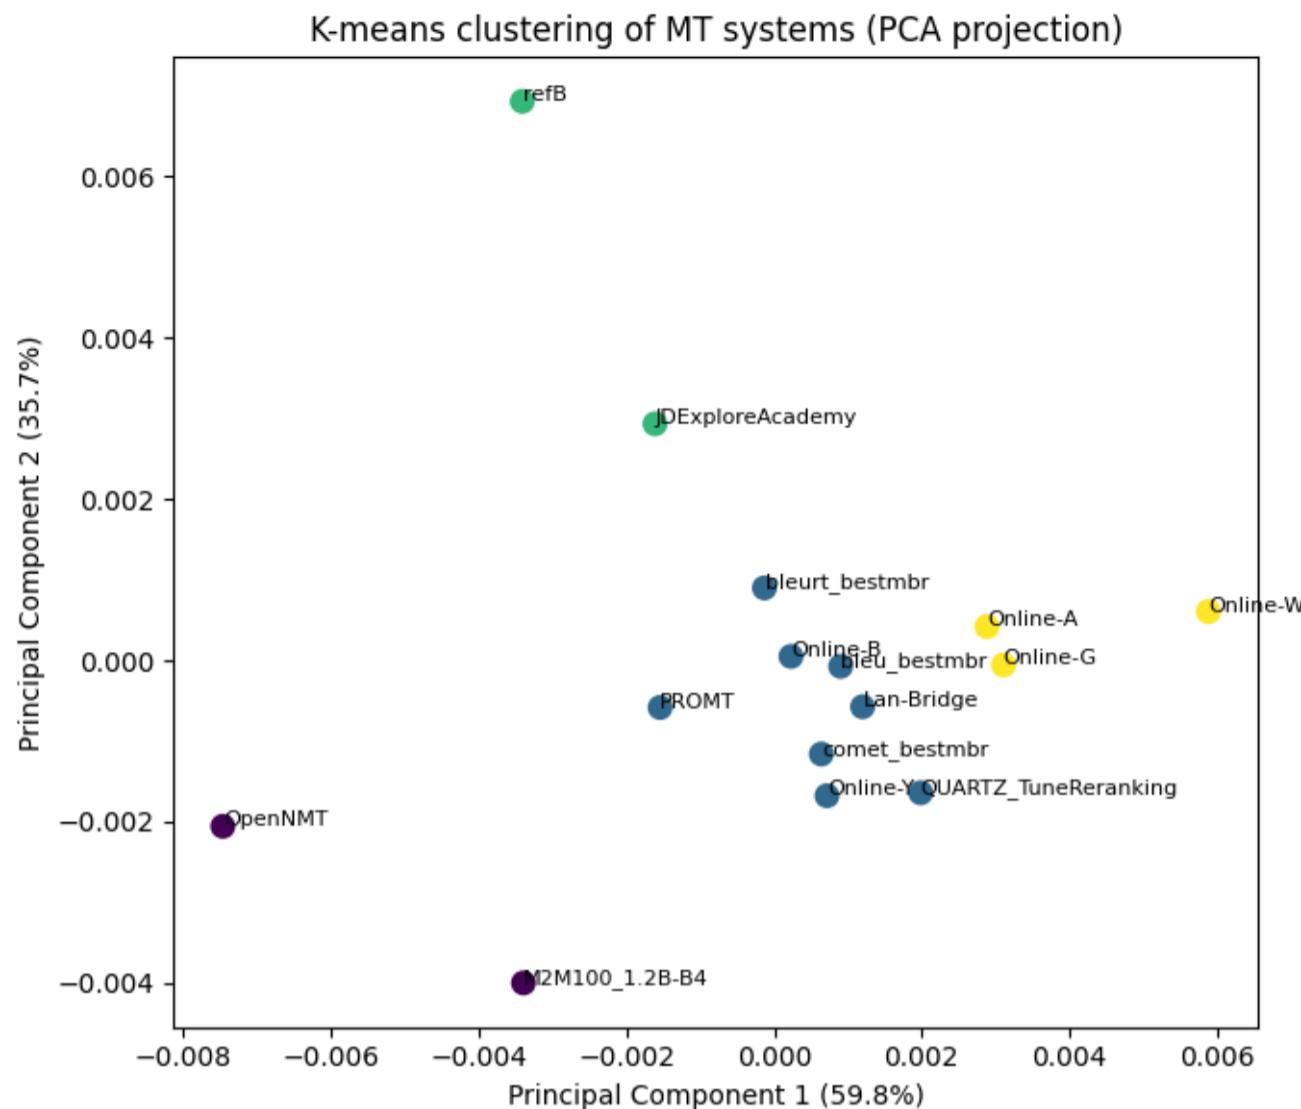
number_of_clusters = 4
kmeans_model = KMeans(
    n_clusters=number_of_clusters,
    n_init=10,
    random_state=0,
)
cluster_labels = kmeans_model.fit_predict(feature_matrix)

pca_model = PCA(n_components=2)
pca_coordinates = pca_model.fit_transform(feature_matrix)

explained_variance = pca_model.explained_variance_ratio_

plt.figure(figsize=(7, 6))
plt.scatter(
    pca_coordinates[:, 0],
    pca_coordinates[:, 1],
    c=cluster_labels,
    s=70,
)
for i, system_name in enumerate(system_names):
    plt.text(
        pca_coordinates[i, 0],
        pca_coordinates[i, 1],
        system_name,
        fontsize=8,
    )

plt.xlabel(f"Principal Component 1 ({explained_variance[0]*100:.1f}%)")
plt.ylabel(f"Principal Component 2 ({explained_variance[1]*100:.1f}%)")
plt.title("K-means clustering of MT systems (PCA projection)")
plt.tight_layout()
plt.show()
```



We arbitrarily chose to cluster with 4 clusters. We tweaked the parameters a bit but 4 gave the most interpretable results. Distances between datapoints in PCA are also interpretable (compared to t-SNE).

First we distinguish a central cluster, of “similar” systems. The yellow cluster groups together the “Online” family (although 2 of its members are in the first cluster). There are also two outlier clusters, M2M and OpenNMT on one hand, refB and JDE on the other hand.

Clustering error profiles groups systems or sentences by how they fail rather than by their overall score. Systems might “behave” the same, while having different architectures or training data. This makes it possible to identify recurring patterns of errors, such as systems that are mostly fluent but inaccurate versus systems that preserve meaning but produce awkward language. Clustering can also help identify outlier systems, that fail in particular cases, as it seems to be the case here.

## 4 MQM and Automatic Metrics

```

import numpy as np
import sacrebleu
import torch

from comet import load_from_checkpoint
from comet.models import download_model

```

```

import pandas as pd

# Define the required columns for metric computation
required_columns = ["system", "segment", "source", "target", "reference"]

# Keep only the columns needed to compute segment/system metrics
metrics_input_dataframe = df[required_columns].copy()

# Drop rows missing any of the text fields needed for metrics
metrics_input_dataframe = metrics_input_dataframe.dropna(subset=["source", "target", "reference"])

# Build the segment-level table (one row per system-segment)
segment_level_metrics_table = metrics_input_dataframe.drop_duplicates(subset=["system", "segment"])

# Build the system-level table skeleton (one row per system)
system_level_metrics_table = (
    segment_level_metrics_table.groupby("system", as_index=False)
    .agg(num_segments=("segment", "nunique"))
)

```

Computing weighted MQM scores following Freitag et al.

```

# Severity weighting scheme
severity_weights = {
    "minor": 1,
    "major": 5,
}

# Map severities to weights
# "no-error" -> NaN (ignored automatically)
df["mqm_weight"] = df["severity"].map(severity_weights)

```

Segment level MQM scores.

```

mqm_segment_scores = df.groupby(["system",
                                "segment"])["mqm_weight"].sum(min_count=1).reset_index(name="mqm_score")
segment_level_metrics_table = segment_level_metrics_table.merge(mqm_segment_scores, on=["system",
                                         "segment"], how="left")

```

System level MQM scores normalized by total tokens per system.

```

df["target_tokens"] = df["target"].str.split().str.len()
tokens_per_system = (df.groupby("system")["target_tokens"].sum())
mqm_system_scores = (df.groupby("system")["mqm_weight"].sum(min_count=1))
mqm_system_scores_norm = mqm_system_scores / tokens_per_system

system_level_metrics_table = system_level_metrics_table.merge(
    mqm_system_scores_norm.reset_index(name="mqm_score_per_token"),
    on="system",
    how="left",
)

```

Computing BLEU and chrF scores at the segment level.

```

bleu_scores = []
chrf_scores = []

for _, row in segment_level_metrics_table.iterrows():
    bleu = sacrebleu.sentence_bleu(row["target"], [row["reference"]]).score
    chrf = sacrebleu.sentence_chrf(row["target"], [row["reference"]]).score
    bleu_scores.append(bleu)
    chrf_scores.append(chrf)

# Add segment-level BLEU and chrF to the segment-level metrics table

```

```
segment_level_metrics_table["BLEU"] = bleu_scores
segment_level_metrics_table["chrF"] = chrf_scores
```

System level BLEU and chrF scores.

```
# Compute corpus-level BLEU and chrF for each system using the segment-level metrics table
system_metric_results = []

for system_name in segment_level_metrics_table["system"].unique():
    system_segments = segment_level_metrics_table[segment_level_metrics_table["system"] == system_name]

    # Collect hypotheses and references for the full system output
    hypotheses = system_segments["target"].astype(str).tolist()
    references = [system_segments["reference"].astype(str).tolist()]

    # Compute corpus-level BLEU (system-level)
    bleu_score = sacrebleu.corpus_bleu(hypotheses, references).score

    # Compute corpus-level chrF (system-level)
    chrf_score = sacrebleu.corpus_chrf(hypotheses, references).score

    system_metric_results.append({
        "system": system_name,
        "BLEU": bleu_score,
        "chrF": chrf_score,
    })

# Put BLEU/chrF results into a separate table (one row per system)
system_level_bleu_chrf_table = pd.DataFrame(system_metric_results)

# Merge BLEU/chrF into the existing system-level metrics table (do NOT overwrite it)
system_level_metrics_table = system_level_metrics_table.merge(system_level_bleu_chrf_table, on="system",
    how="left")
```

COMET score.

```
# Load the COMET-DA model (reference-based, suitable for MQM correlation)
model_path = download_model("Unbabel/wmt22-comet-da")
comet_model = load_from_checkpoint(model_path)

# Prepare COMET inputs from the segment-level metrics table (one row per system-segment)
comet_inputs = [
    {
        "src": row["source"],
        "mt": row["target"],
        "ref": row["reference"],
    }
    for _, row in segment_level_metrics_table.iterrows()
]

# Run COMET to obtain segment-level scores
comet_output = comet_model.predict(
    comet_inputs,
    batch_size=32,
    gpus=1 if torch.cuda.is_available() else 0,
    num_workers=1,
)

# Add segment-level COMET scores to the segment-level metrics table
segment_level_metrics_table["COMET"] = comet_output["scores"]

# Aggregate system-level COMET as the mean of segment-level scores
comet_system_scores = segment_level_metrics_table.groupby("system")["COMET"].mean()

# Merge system-level COMET into the existing system-level metrics table
system_level_metrics_table = system_level_metrics_table.merge(
    comet_system_scores.rename("COMET"),
```

```

    on="system",
    how="left",
)

segment_level_metrics_table = pd.read_csv("data/segment_level_metrics.csv")
system_level_metrics_table = pd.read_csv("data/system_level_metrics.csv")

display(segment_level_metrics_table)
display(system_level_metrics_table)

```

	system	segment	source	target	reference
0	JDExploreAcademy	1	Avalanche at Was...	Lawine in Skigeb...	Lawine in Ski-Re...
1	Lan-Bridge	1	Avalanche at Was...	Lawine im Skigeb...	Lawine in Ski-Re...
2	M2M100_1.2B-B4	1	Avalanche at Was...	<v>Avalanche</v>...	Lawine in Ski-Re...
3	Online-A	1	Avalanche at Was...	Lawine im Skigeb...	Lawine in Ski-Re...
4	Online-B	1	Avalanche at Was...	Lawine im Skigeb...	Lawine in Ski-Re...
...	...	...	...	...	...
145	QUARTZ_TuneReran...	10	"Skiers can trav...	"Im Staatswald k...	„Skifahrer könne...
146	bleu_bestmbr	10	"Skiers can trav...	„Skifahrer könne...	„Skifahrer könne...
147	bleurt_bestmbr	10	"Skiers can trav...	„Skifahrer könne...	„Skifahrer könne...
148	comet_bestmbr	10	"Skiers can trav...	„Skifahrer könne...	„Skifahrer könne...
149	refB	10	"Skiers can trav...	„Skifahrer könne...	„Skifahrer könne...

	system	num_segments	mqm_score_per_token	BLEU	chrF	COMET
0	JDExploreAcademy	10	0.054599	32.666188	61.775083	0.822923
1	Lan-Bridge	10	0.064303	24.708933	59.038817	0.761160
2	M2M100_1.2B-B4	10	0.091230	21.711807	55.987241	0.721010
3	Online-A	10	0.060145	27.270501	59.721487	0.786887
4	Online-B	10	0.054561	25.004807	59.665942	0.770541
5	Online-G	10	0.062263	26.326593	60.875855	0.758391
6	Online-W	10	0.044980	34.177264	62.279745	0.845575
7	Online-Y	10	0.062435	25.471897	59.843116	0.740286
8	OpenNMT	10	0.075304	27.706626	60.139061	0.754235
9	PROMT	10	0.070187	25.850586	61.062388	0.763773
10	QUARTZ_TuneReran...	10	0.065215	16.270910	51.311063	0.793082
11	bleu_bestmbr	10	0.054473	28.446967	62.207114	0.784587
12	bleurt_bestmbr	10	0.055316	24.415303	58.161545	0.818754
13	comet_bestmbr	10	0.058619	26.081795	59.319196	0.830675
14	refB	10	0.047307	25.122366	59.532564	0.811405

System level correlations.

```

from scipy.stats import pearsonr, spearmanr

system_correlation_results = []

for metric_name in ["BLEU", "chrF", "COMET"]:

```

```

pearson_r, _ = pearsonr(
    system_level_metrics_table["mqm_score_per_token"],
    system_level_metrics_table[metric_name],
)
spearman_r, _ = spearmanr(
    system_level_metrics_table["mqm_score_per_token"],
    system_level_metrics_table[metric_name],
)

system_correlation_results.append({
    "metric": metric_name,
    "pearson_r": pearson_r,
    "spearman_r": spearman_r,
})

system_correlation_table = pd.DataFrame(system_correlation_results)

system_correlation_table

```

	metric	pearson_r	spearman_r
0	BLEU	-0.474604	-0.417857
1	chrF	-0.420903	-0.392857
2	COMET	-0.778326	-0.707143

The system level correlations are negative, as expected. There is very subtle difference between Pearson and Spearman correlations, so interpreting between the two is not very meaningful here.

We notice that there is a stronger correlation between MQM and COMET than between MQM and BLEU or chrF. We will look at the distribution of the data more closely in the scatter plots.

Segment level correlations.

```

segment_correlation_results = []

for metric_name in ["BLEU", "chrF", "COMET"]:
    pearson_r, _ = pearsonr(
        segment_level_metrics_table["mqm_score"],
        segment_level_metrics_table[metric_name],
    )
    spearman_r, _ = spearmanr(
        segment_level_metrics_table["mqm_score"],
        segment_level_metrics_table[metric_name],
    )

    segment_correlation_results.append({
        "metric": metric_name,
        "pearson_r": pearson_r,
        "spearman_r": spearman_r,
    })

segment_correlation_table = pd.DataFrame(segment_correlation_results)

segment_correlation_table

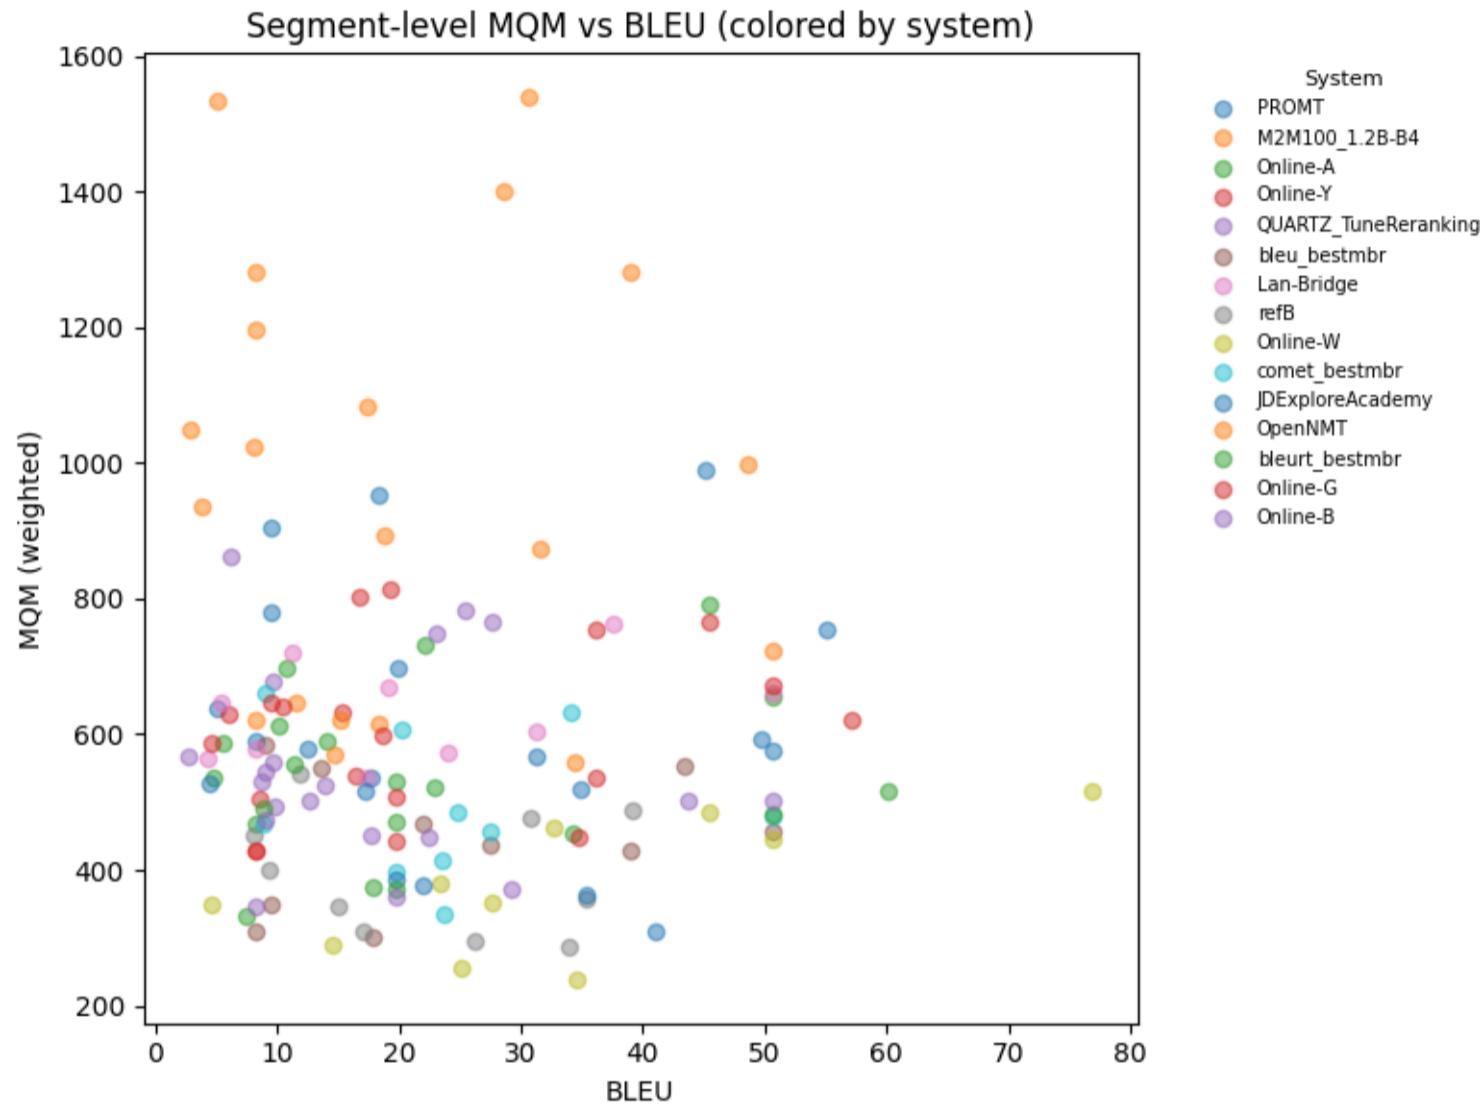
```

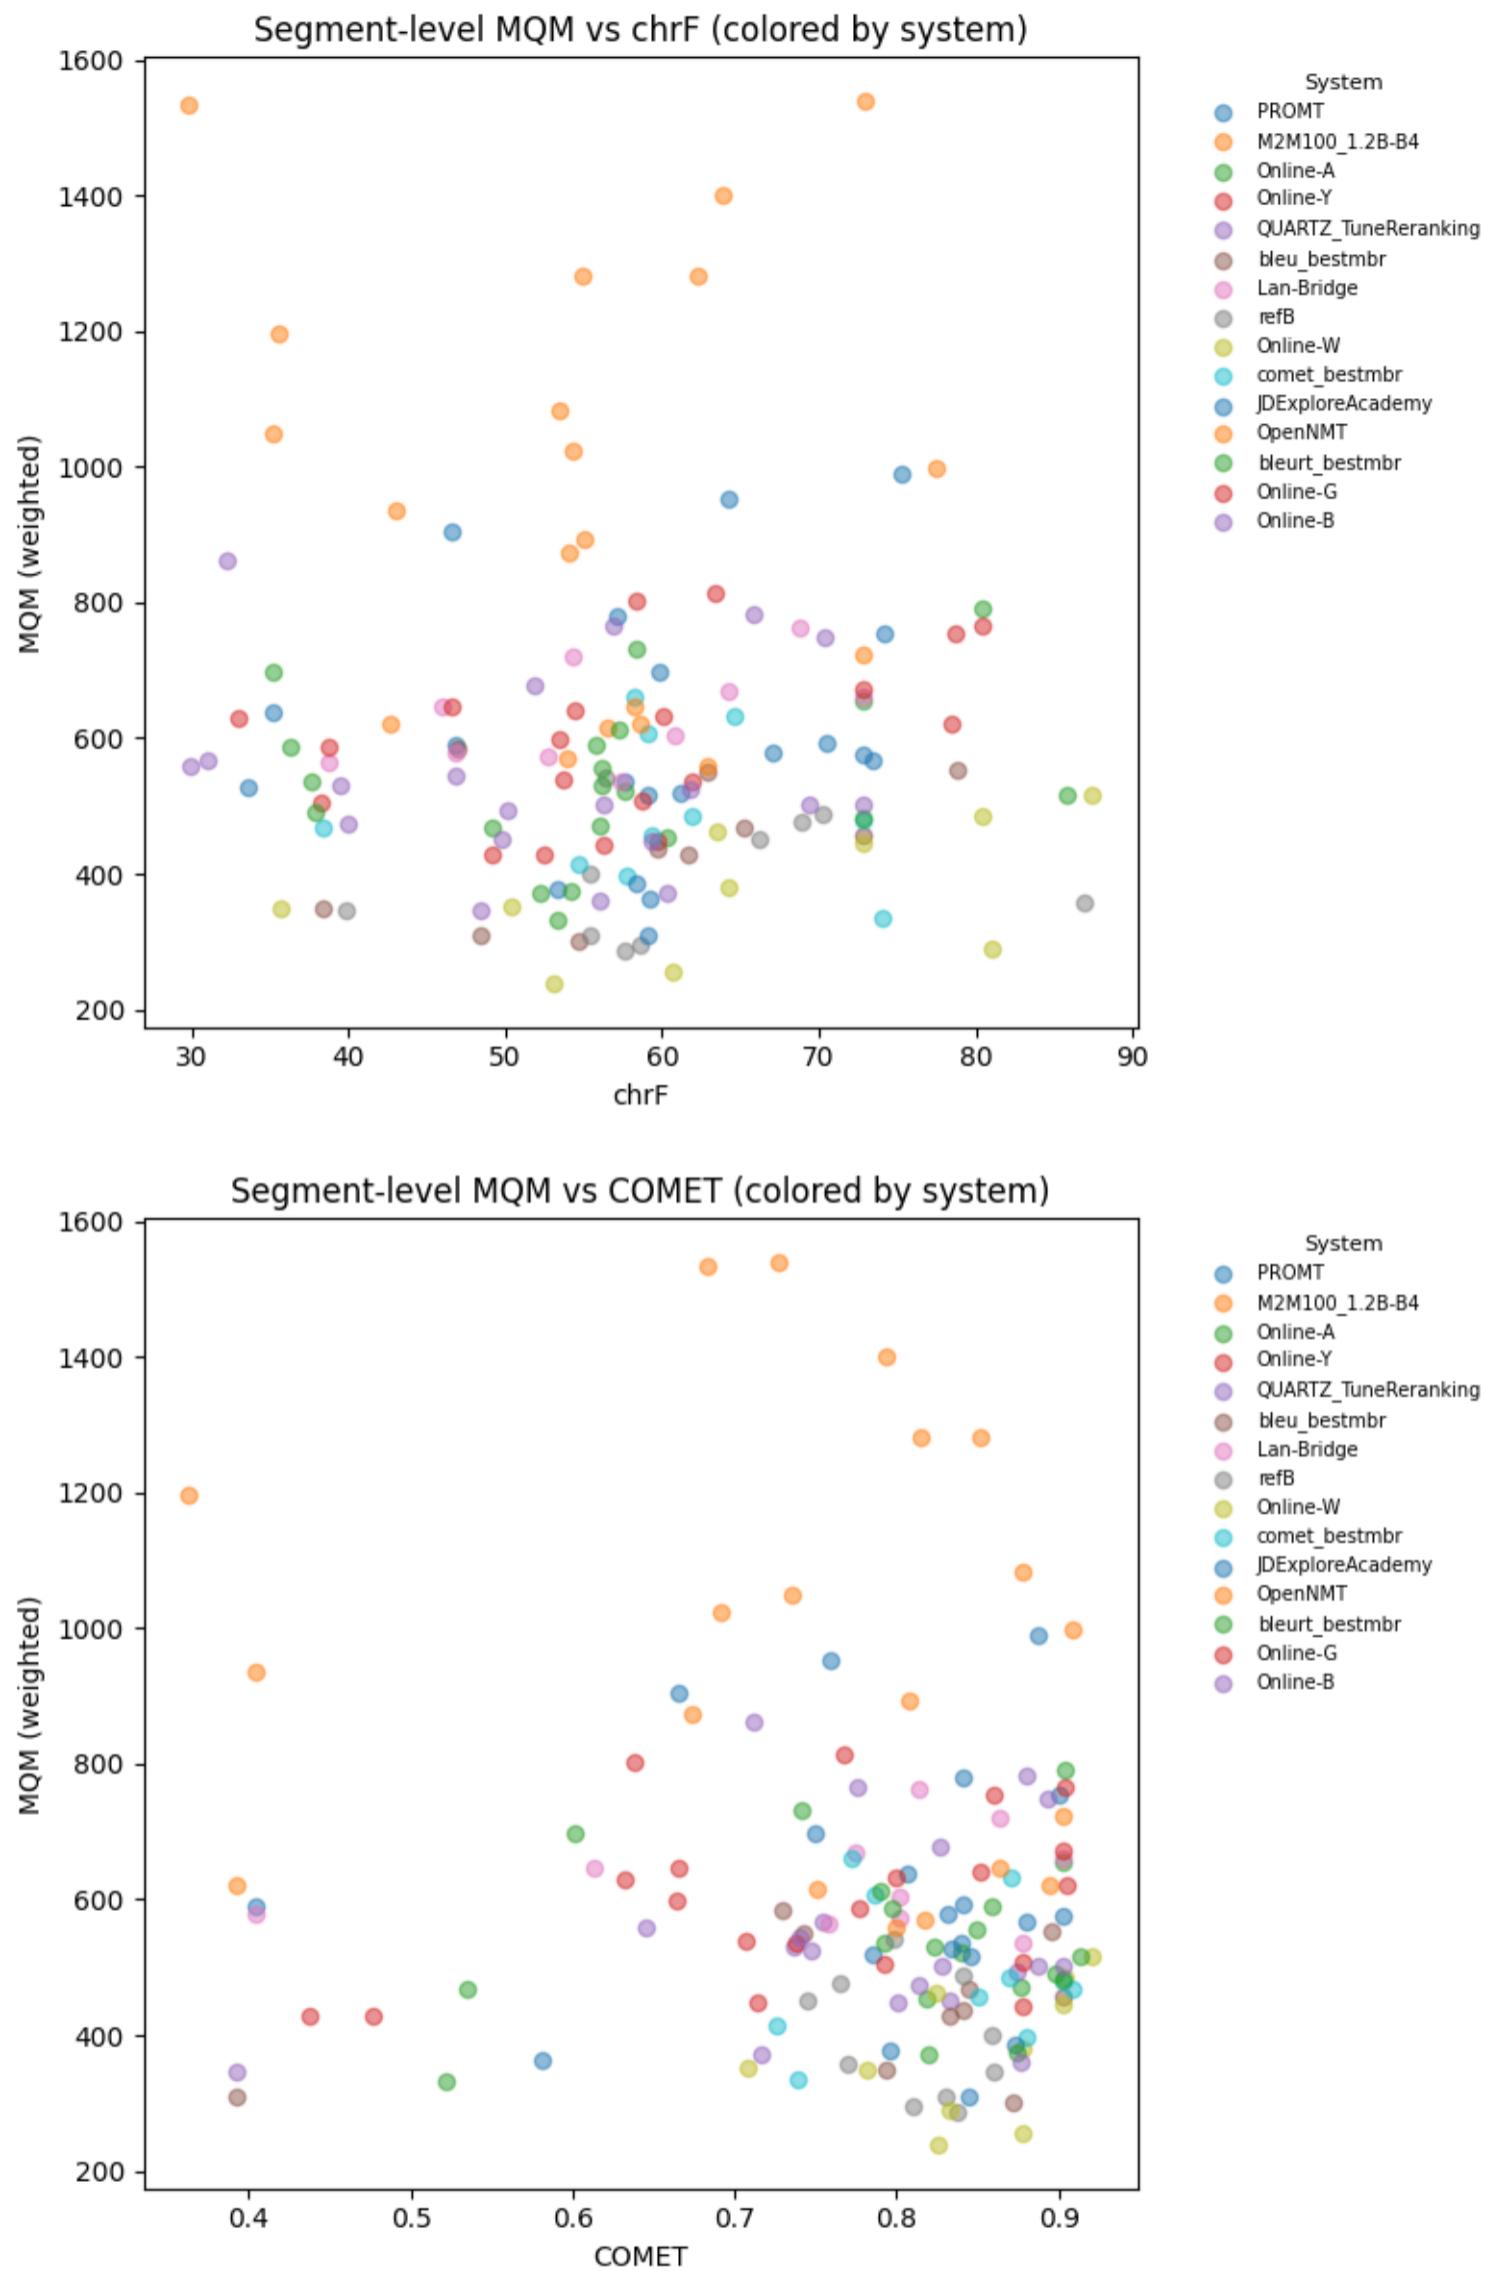
	metric	pearson_r	spearman_r
0	BLEU	-0.030795	-0.073729
1	chrF	-0.024369	0.032225

metric	pearson_r	spearman_r
2 COMET	-0.088716	-0.071358

Although the correlations maintain their negative sign, we can see sharp decrease. This is also expected, as automatic metrics are designed to be computed at the system level, especially for BLEU and chrF. Segment level correlations are much more noisier.

Plotting segment level MQM vs automatic metrics.



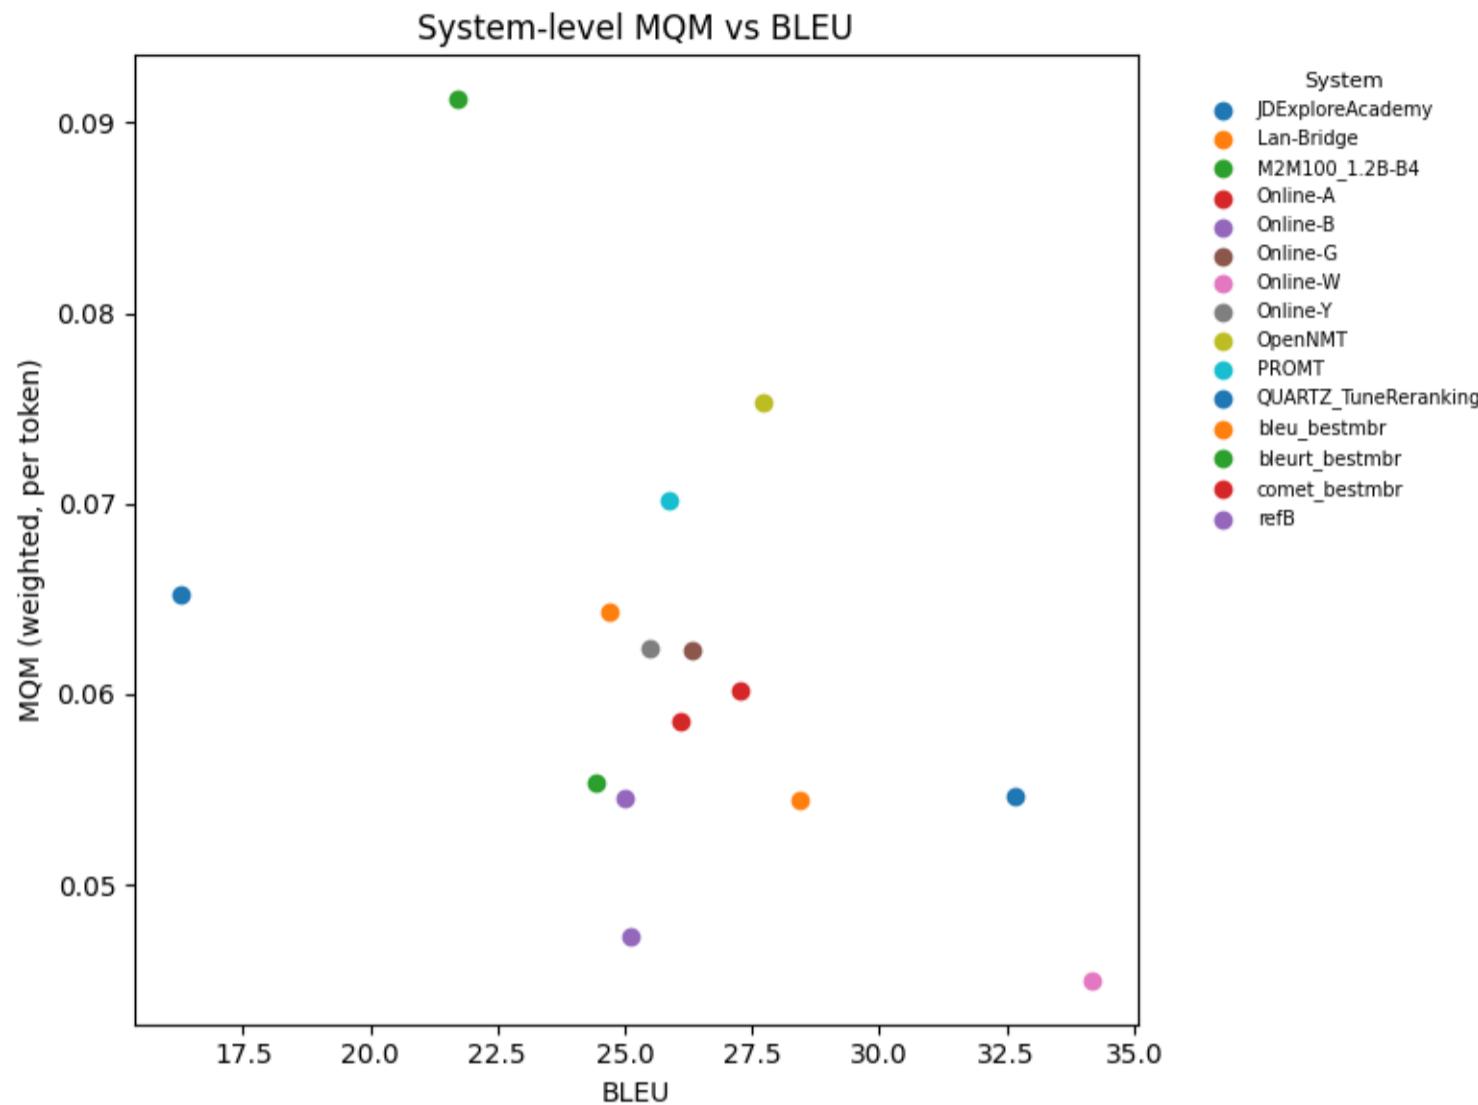


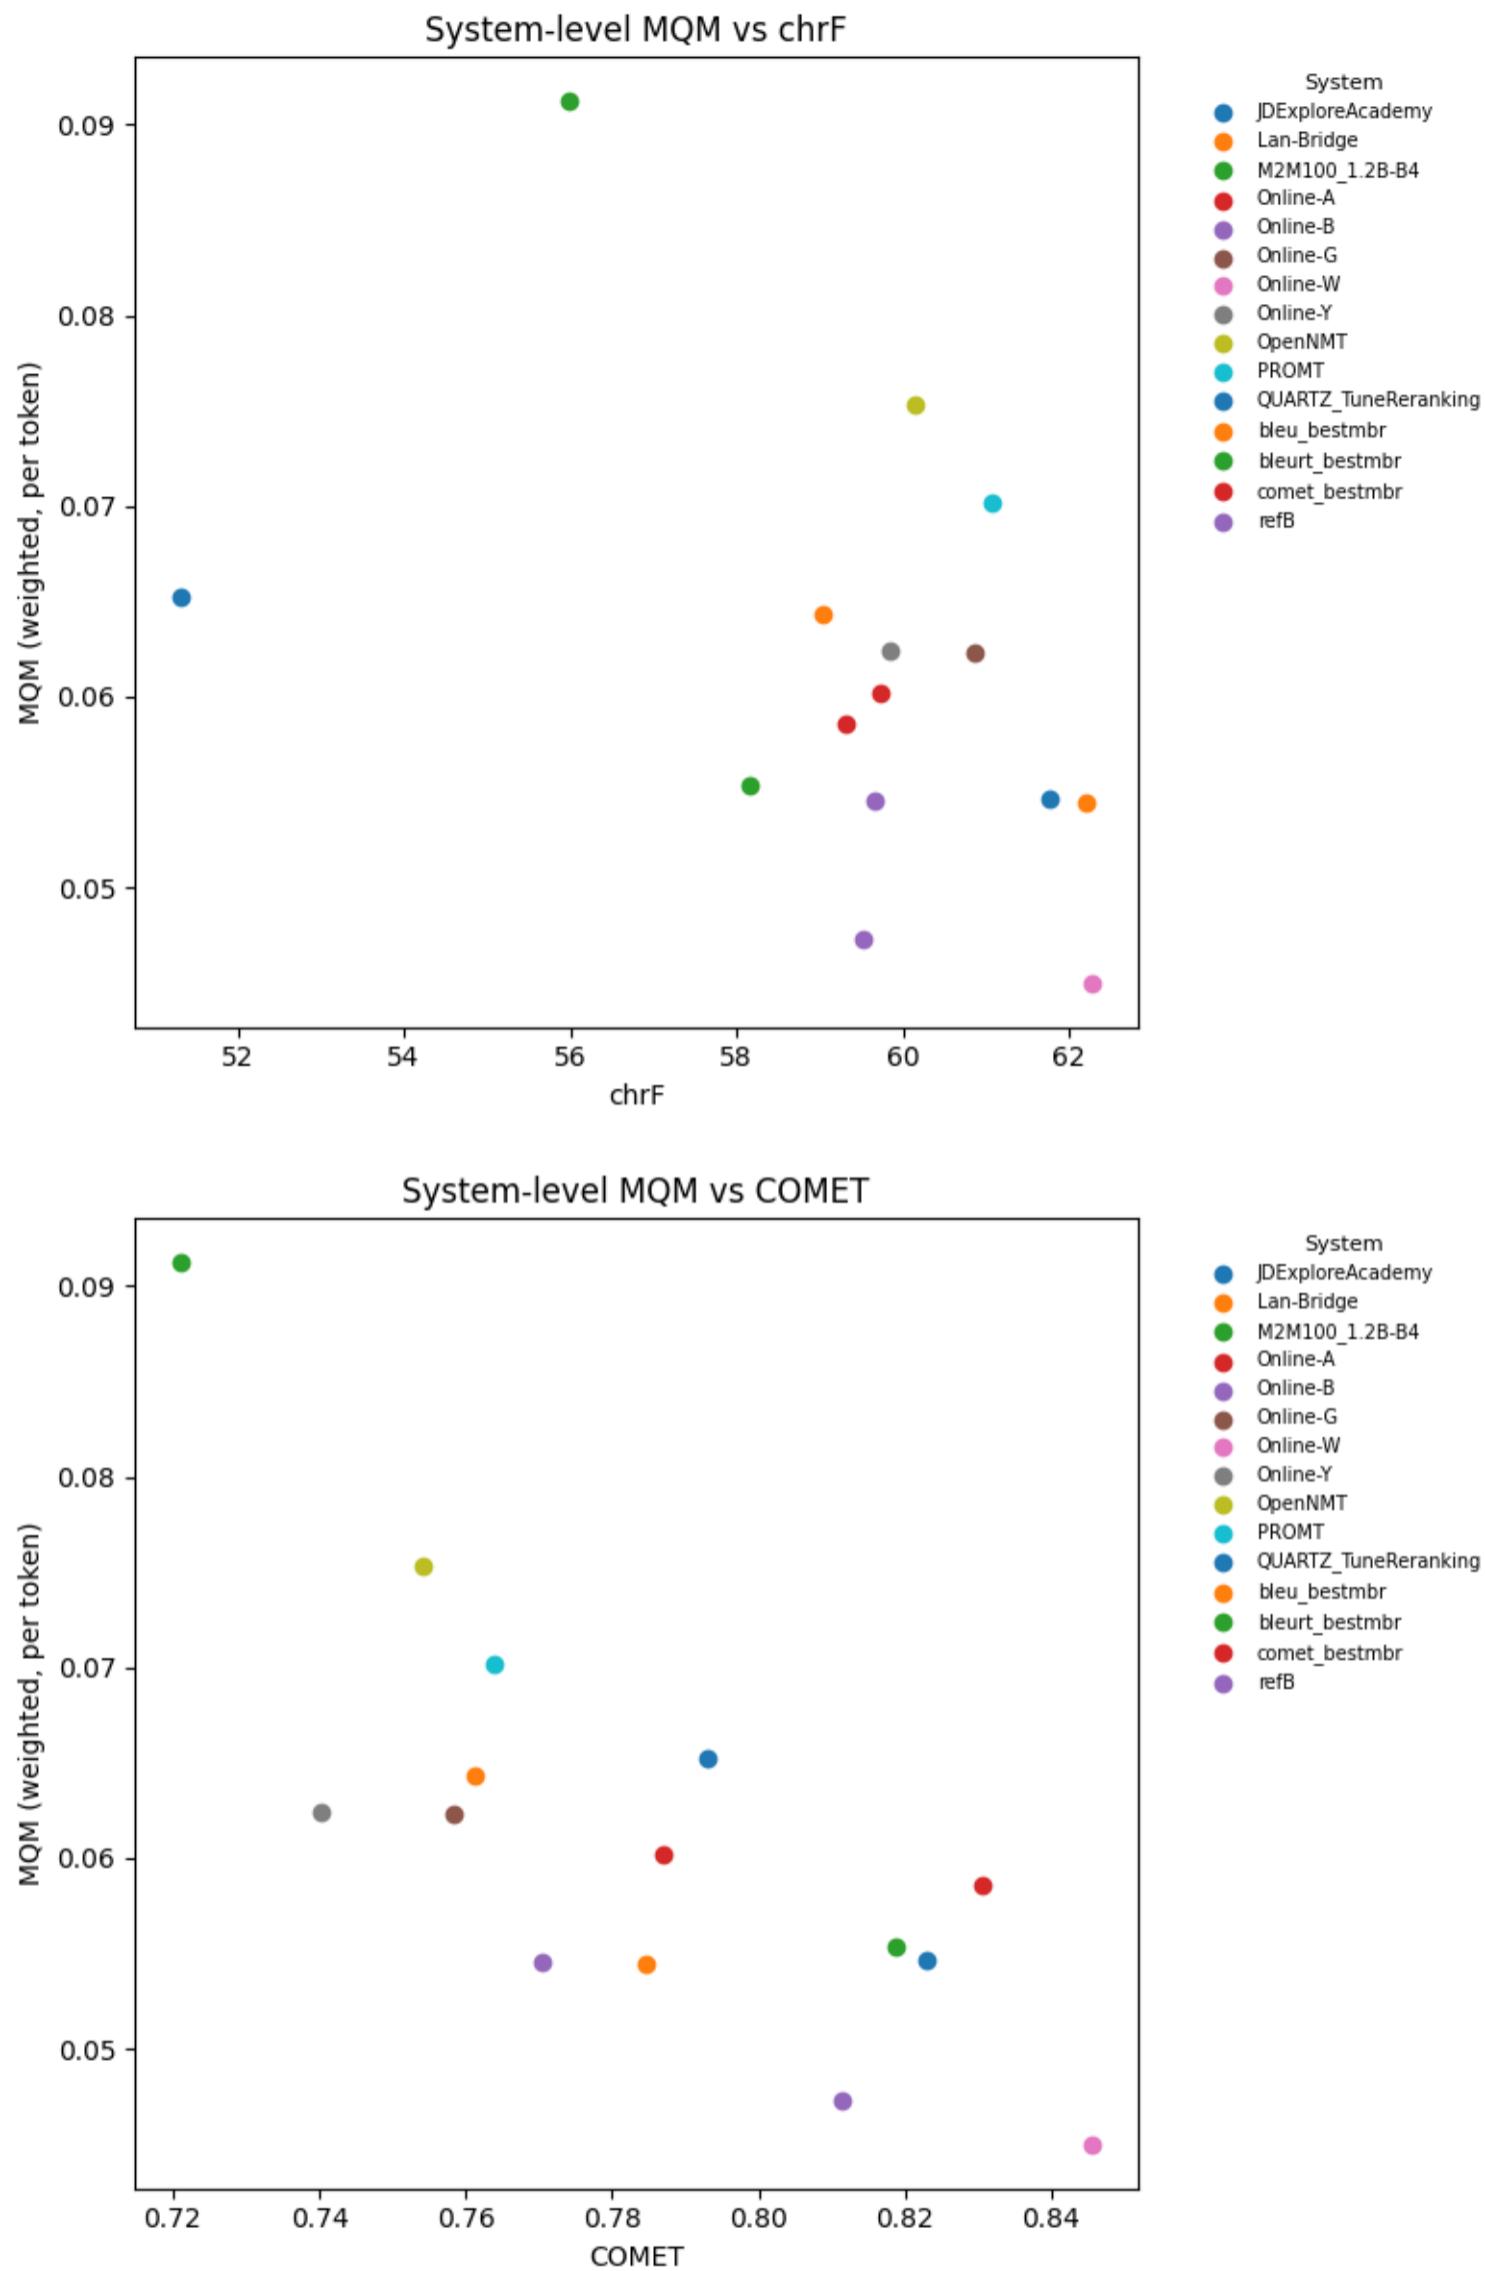
Grouping the datapoints between systems, we do not see a clear pattern, if not for M2M

that is terrible...

The scatter plot between MQM and COMET shows the clearer negative correlation pattern, with higher COMET scores corresponding to lower MQM error scores. BLEU and chrF show much more noise, with many segments average chrF scores and low BLEU scores corresponding to a wide range of MQM scores.

And at system level.





At the system level, the correlations become much clearer. COMET reveals much more

about system level quality, as datapoints are more spread out along both axes. BLEU and chrF also show the expected negative correlation, but systems are slightly more clustered together on the x axis. Overall COMET seems to be a better predictor of MQM scores than BLEU or chrF, both at segment and system level.

## 5 Linguistic factors and correlation

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.cluster import KMeans
```

Computing surface features:

- Character length of the source sentence
- Hedran's C (a normalized TTR).

The formula for Hedran's C is:

$$C = \frac{\log V}{\log N}$$

where V is the number of unique word types in the segment and N is the total number of tokens.

```
# Character length of the source sentence
segment_level_metrics_table["char_length"] = segment_level_metrics_table["source"].astype(str).str.len()

# Tokenize source sentences (simple whitespace tokenization)
source_tokens = segment_level_metrics_table["source"].astype(str).str.split()

# Number of tokens per segment
segment_level_metrics_table["num_tokens"] = source_tokens.str.len()

# Number of unique tokens per segment
segment_level_metrics_table["num_types"] = source_tokens.apply(lambda toks: len(set(toks)) if toks else 0)

# Herdan's C: log(V) / log(N); undefined for N <= 1 → set to NaN
segment_level_metrics_table["herdan_c"] = np.where(
    segment_level_metrics_table["num_tokens"] > 1,
    np.log(segment_level_metrics_table["num_types"]) / np.log(segment_level_metrics_table["num_tokens"]),
    np.nan,
)
```

TF-IDF clustering of source sentences.

```
# Build TF-IDF vectors from source sentences
tfidf_vectorizer = TfidfVectorizer(
    lowercase=True,
    max_features=5000,
)

tfidf_matrix = tfidf_vectorizer.fit_transform(
    segment_level_metrics_table["source"].astype(str)
)

# Number of clusters (keep small and interpretable)
num_clusters = 5
```

```

kmeans_model = KMeans(
    n_clusters=num_clusters,
    n_init=10,
    random_state=0,
)

cluster_ids = kmeans_model.fit_predict(tfidf_matrix)

# Append cluster labels to the segment-level metrics table
segment_level_metrics_table["cluster_id"] = cluster_ids

```

```

import pandas as pd

# Compute descriptive statistics per TF-IDF cluster
cluster_statistics = (
    segment_level_metrics_table
        .groupby("cluster_id")
        .agg(
            num_segments=("segment", "count"),
            avg_char_length=("char_length", "mean"),
            std_char_length=("char_length", "std"),
            avg_herdan_c=("herdan_c", "mean"),
            std_herdan_c=("herdan_c", "std"),
            avg_mqm=("mqm_score", "mean"),
            avg_bleu=("BLEU", "mean"),
            avg_chrf=("chrF", "mean"),
            avg_comet=("COMET", "mean"),
        )
        .reset_index()
)

```

cluster\_statistics

	cluster_id	num_segments	avg_char_length	std_char_length	avg_herdan_c	std_herdan_c	avg
0	0	45	188.333333	114.347954	0.978628	0.022871	634.
1	1	30	152.000000	34.581239	0.983601	0.002260	627.
2	2	15	35.000000	0.000000	1.000000	0.000000	480.
3	3	30	207.233333	9.474770	0.972140	0.010851	637.
4	4	30	130.000000	62.042811	0.983785	0.016492	503.

Overall, for segment level computations, Hedran's C is not very informative, as most tokens are also counted as types (well if computed per sentence, this is trivial). On the other hand, the clusters vary based on the average character length and standard deviation.

- Cluster 1 has high character length and a high standard deviation.
- Cluster 2 has high character length and low standard deviation.
- Cluster 3 has low character length and a low standard deviation (sd of zero, indicating that the segments are of the same length).
- Cluster 4 has the highest character length and a very low standard deviation.
- Cluster 4 is relatively similar to cluster 2.

Recomputing correlations per cluster.

```

import pandas as pd
from scipy.stats import pearsonr, spearmanr

```

```

# Collect TF-IDF cluster-level correlation results in a structured table
cluster_correlation_results = []

for cluster_id in sorted(segment_level_metrics_table["cluster_id"].unique()):
    cluster_data = segment_level_metrics_table[
        segment_level_metrics_table["cluster_id"] == cluster_id
    ]

    for metric_name in ["BLEU", "chrF", "COMET"]:
        pearson_r, _ = pearsonr(
            cluster_data["mqm_score"],
            cluster_data[metric_name],
        )
        spearman_r, _ = spearmanr(
            cluster_data["mqm_score"],
            cluster_data[metric_name],
        )

        cluster_correlation_results.append({
            "cluster_id": cluster_id,
            "metric": metric_name,
            "num_segments": len(cluster_data),
            "pearson_r": pearson_r,
            "spearman_r": spearman_r,
        })

# Convert to DataFrame
tfidf_cluster_correlation_table = pd.DataFrame(cluster_correlation_results)

tfidf_cluster_correlation_table

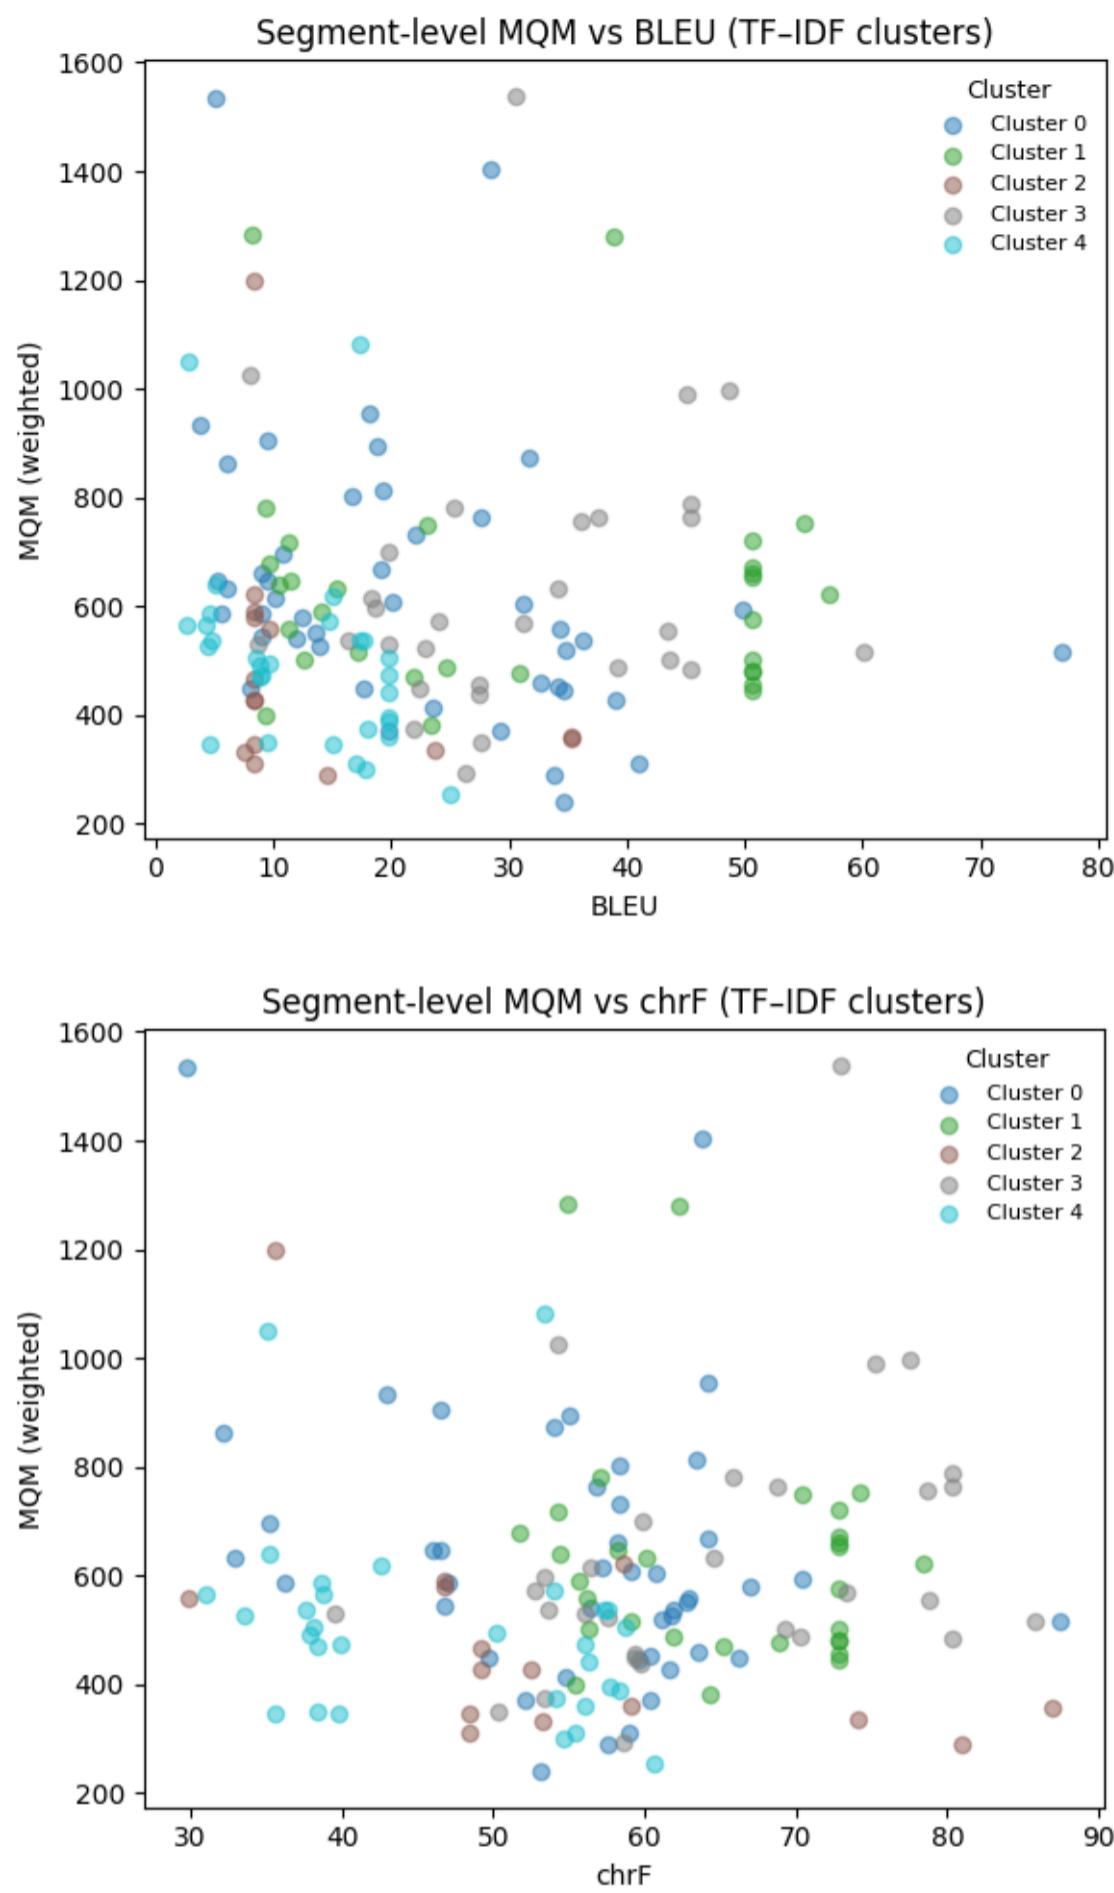
```

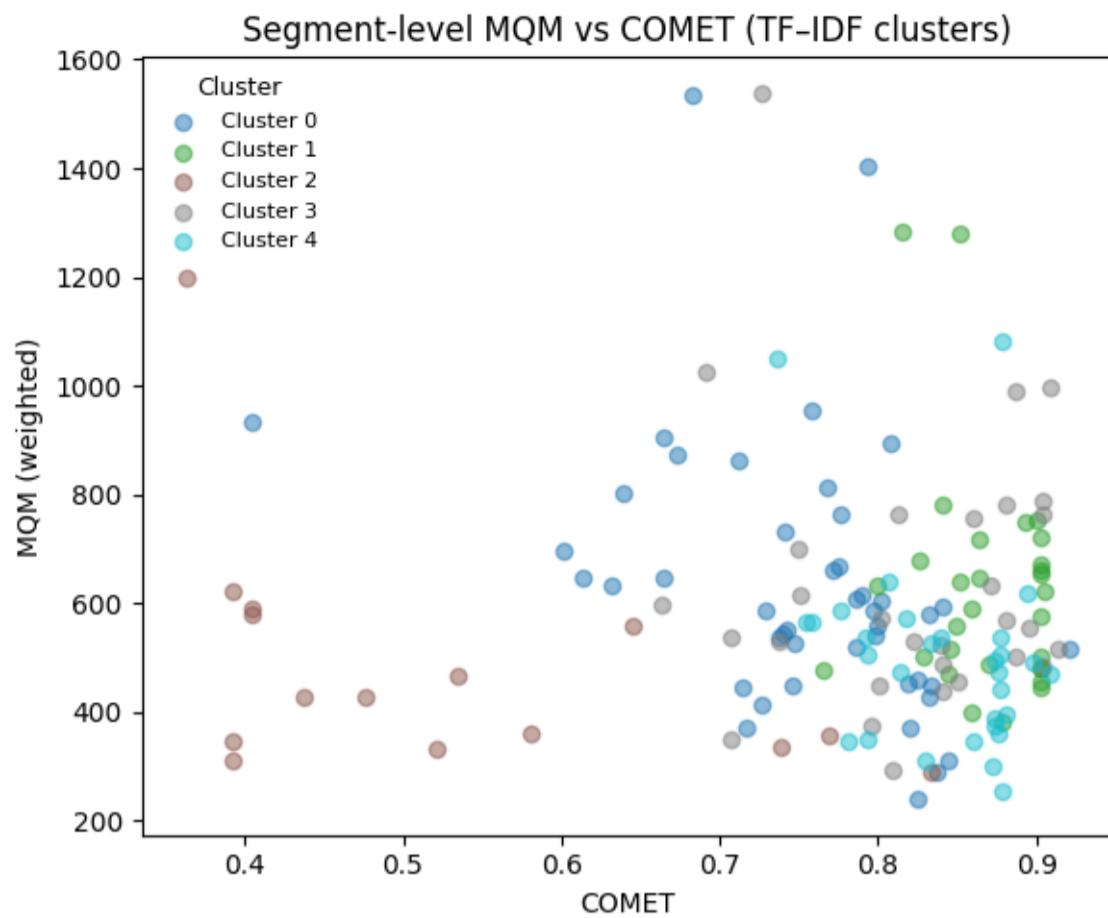
	cluster_id	metric	num_segments	pearson_r	spearman_r
0	0	BLEU	45	-0.332387	-0.480482
1	0	chrF	45	-0.306500	-0.222808
2	0	COMET	45	-0.387375	-0.491600
3	1	BLEU	30	-0.105968	-0.140210
4	1	chrF	30	-0.197769	-0.134996
5	1	COMET	30	-0.177681	-0.120941
6	2	BLEU	15	-0.310248	-0.223870
7	2	chrF	15	-0.535192	-0.554563
8	2	COMET	15	-0.465373	-0.500895
9	3	BLEU	30	0.124528	0.132888
10	3	chrF	30	0.344018	0.335856
11	3	COMET	30	-0.057990	0.091445
12	4	BLEU	30	-0.350920	-0.477662
13	4	chrF	30	-0.264445	-0.367824
14	4	COMET	30	-0.292790	-0.235648

```

/tmp/ipython-input-1955519037.py:5: MatplotlibDeprecationWarning: The get_cmap function was deprecated in
↔ Matplotlib 3.7 and will be removed in 3.11. Use ``matplotlib.colormaps[name]`` or
↔ ``matplotlib.colormaps.get_cmap()`` or ``pyplot.get_cmap()`` instead.
color_map = plt.cm.get_cmap("tab10", len(cluster_ids))

```





The scatter plots are much more telling than the table of correlations.

Cluster 2 is clearly isolated in the COMET vs MQM scatter plot; it corresponds to the cluster with the shorter segments, and confirms that automatic metrics are less reliable on short segments. Indeed this cluster has low MQM scores but is affected very low COMET scores.

As at the segment level BLEU and chrF are much more noisy, we are not able to distinguish a clear separation between our clusters.

We can also try with embeddings-based clustering.

```
from sentence_transformers import SentenceTransformer

# Load a lightweight sentence embedding model
embedding_model = SentenceTransformer("all-MiniLM-L6-v2")

# Encode source sentences into dense vectors
source_sentences = segment_level_metrics_table["source"].astype(str).tolist()
sentence_embeddings = embedding_model.encode(
    source_sentences,
    batch_size=32,
    show_progress_bar=True,
)

# Cluster sentences in embedding space
num_clusters = 5
kmeans_model = KMeans(
    n_clusters=num_clusters,
    n_init=10,
    random_state=0,
)
embedding_cluster_ids = kmeans_model.fit_predict(sentence_embeddings)

# Append embedding-based cluster IDs to the segment-level table
segment_level_metrics_table["embedding_cluster"] = embedding_cluster_ids
```

```

embedding_cluster_correlation_results = []

for cluster_id in sorted(segment_level_metrics_table["embedding_cluster"].unique()):
    cluster_data = segment_level_metrics_table[
        segment_level_metrics_table["embedding_cluster"] == cluster_id
    ]

    for metric_name in ["BLEU", "chrF", "COMET"]:
        pearson_r, _ = pearsonr(
            cluster_data["mqm_score"],
            cluster_data[metric_name],
        )
        spearman_r, _ = spearmanr(
            cluster_data["mqm_score"],
            cluster_data[metric_name],
        )

        embedding_cluster_correlation_results.append({
            "embedding_cluster": cluster_id,
            "metric": metric_name,
            "num_segments": len(cluster_data),
            "pearson_r": pearson_r,
            "spearman_r": spearman_r,
        })

# Convert to DataFrame
embedding_cluster_correlation_table = pd.DataFrame(embedding_cluster_correlation_results)

embedding_cluster_correlation_table

```

Batches: 0% | 0/5 [00:00<?, ?it/s]

		embedding_cluster	metric	num_segments	pearson_r	spearman_r
0	0		BLEU	30	-0.350920	-0.477662
1	0		chrF	30	-0.264445	-0.367824
2	0		COMET	30	-0.292790	-0.235648
3	1		BLEU	30	-0.132901	-0.090565
4	1		chrF	30	-0.315341	-0.346646
5	1		COMET	30	-0.260122	-0.369785
6	2		BLEU	15	-0.016728	-0.025000
7	2		chrF	15	0.123890	0.425000
8	2		COMET	15	-0.571589	-0.496429
9	3		BLEU	30	-0.105968	-0.140210
10	3		chrF	30	-0.197769	-0.134996
11	3		COMET	30	-0.177681	-0.120941
12	4		BLEU	45	0.182905	0.209236
13	4		chrF	45	0.121521	0.149173
14	4		COMET	45	0.104267	0.216490

```

/tmp/ipython-input-1973396413.py:5: MatplotlibDeprecationWarning: The get_cmap function was deprecated in
↔ Matplotlib 3.7 and will be removed in 3.11. Use ``matplotlib.colormaps[name]`` or
↔ ``matplotlib.colormaps.get_cmap()`` or ``pyplot.get_cmap()`` instead.
color_map = plt.cm.get_cmap("tab10", len(cluster_ids))

```

